

## **STUDENT 8: AGRICULTURAL PRODUCTION & SUPPLY**

**A1**

```
41 -- Test connection to Node_B
42 SELECT * FROM dblink(
43     'host=localhost port=5432 dbname=Node_B user=postgres password=Bobo1999@',
44     'SELECT field_id, field_name FROM Field'
45 ) AS remote_field(field_id INTEGER, field_name VARCHAR(100));
```

Data Output Messages Notifications

	field_id	field_name
1	1	North Field
2	2	South Field

Showing rows: 1 to 2 Page No: 1 of 1

```
55 -- Insert 5 rows into Harvest_A (on Node_A)
56 INSERT INTO Harvest_A VALUES (1, 1, 101, '2023-06-15', 500);
57 INSERT INTO Harvest_A VALUES (2, 1, 101, '2023-06-20', 450);
58 INSERT INTO Harvest_A VALUES (3, 2, 102, '2023-07-10', 300);
59 INSERT INTO Harvest_A VALUES (4, 2, 102, '2023-07-12', 320);
60 INSERT INTO Harvest_A VALUES (5, 1, 101, '2023-08-01', 480);
61
62 -- Verify the data was inserted
63 SELECT * FROM Harvest_A;
64
```

Data Output Messages Notifications

	harvest_id	field_id	crop_id	harvest_date	yield_kg
1	1	1	101	2023-06-15	500
2	2	1	101	2023-06-20	450
3	3	2	102	2023-07-10	300
4	4	2	102	2023-07-12	320
5	5	1	101	2023-08-01	480

Showing rows: 1 to 5 Page No: 1

```

52 -- Insert 5 rows into Harvest_B (on Node_B via dblink)
53 INSERT INTO Harvest_B VALUES (6, 1, 102, '2023-09-05', 200);
54 INSERT INTO Harvest_B VALUES (7, 2, 101, '2023-09-10', 400);
55 INSERT INTO Harvest_B VALUES (8, 2, 101, '2023-09-15', 410);
56 INSERT INTO Harvest_B VALUES (9, 1, 102, '2023-10-01', 210);
57 INSERT INTO Harvest_B VALUES (10, 2, 102, '2023-10-05', 220);
58
59 -- Verify the data was inserted
60 SELECT * FROM Harvest_B;
61

```

Data Output Messages Notifications

	harvest_id [PK] integer	field_id integer	crop_id integer	harvest_date date	yield_kg numeric
1	6	1	102	2023-09-05	200
2	7	2	101	2023-09-10	400
3	8	2	101	2023-09-15	410
4	9	1	102	2023-10-01	210
5	10	2	102	2023-10-05	220

```

85 -- 4. Validate the fragmentation and view.
86 -- Count should match: (5) + (5) = 10
87 SELECT 'Harvest_A' AS fragment, COUNT(*) FROM Harvest_A
88 UNION ALL
89 SELECT 'Harvest_B' AS fragment, COUNT(*) FROM dblink(
90   'host=localhost port=5432 dbname=Node_B user=postgres password=Bobo1999@',
91   'SELECT COUNT(*) FROM Harvest_B'
92 ) AS remote_count(count BIGINT)
93 UNION ALL
94 SELECT 'Harvest_ALLL' AS fragment, COUNT(*) FROM Harvest_ALLL;
95

```

Data Output Messages Notifications

	fragment text	count bigint
1	Harvest_A	5
2	Harvest_B	1
3	Harvest_AL...	10

```

92    -- Count should match: (10) + (5) = 15
93    SELECT 'Harvest_A' AS fragment, COUNT(*) FROM Harvest_A
94    UNION ALL
95    SELECT * FROM dblink(
96        'host=localhost port=5432 dbname=Node_B user=postgres password=Bobo1999@',
97        'SELECT ''Harvest_B'' AS fragment, COUNT(*) FROM Harvest_B'
98    ) AS remote_count(fragment TEXT, count BIGINT)
99    UNION ALL
100   SELECT 'Harvest_ALLLL' AS fragment, COUNT(*) FROM Harvest_ALLLL;
101

```

Data Output Messages Notifications

Showing rows: 1 to 3 Page No: 1

	fragment text	count bigint
1	Harvest_A	10
2	Harvest_B	5
3	Harvest_ALL...	15

```

103    -- Checksum validation using a simple MOD on primary key.
104    SELECT 'Harvest_A' AS fragment, SUM(MOD(harvest_id, 97)) AS checksum FROM Harvest_A
105    UNION ALL
106    SELECT * FROM dblink(
107        'host=localhost port=5432 dbname=Node_B user=postgres password=Bobo1999@',
108        'SELECT ''Harvest_B'' AS fragment, SUM(MOD(harvest_id, 97)) FROM Harvest_B'
109    ) AS remote_checksum(fragment TEXT, checksum NUMERIC)
110    UNION ALL
111    SELECT 'Harvest_ALLLL' AS fragment, SUM(MOD(harvest_id, 97)) FROM Harvest_ALLLL;

```

Data Output Messages Notifications

Showing rows: 1 to 3 Page No: 1 of 1

	fragment text	checksum numeric
1	Harvest_A	30
2	Harvest_B	40
3	Harvest_ALL...	70

```
114 -- Test the view works
115 SELECT * FROM Harvest_ALLLL ORDER BY harvest_id;
116
```

Data Output Messages Notifications

Showing rows: 1 to 15

	harvest_id integer	field_id integer	crop_id integer	harvest_date date	yield_kg numeric
1	1	1	101	2023-06-15	500
2	1	1	101	2023-06-15	500
3	2	1	101	2023-06-20	450
4	2	1	101	2023-06-20	450
5	3	2	102	2023-07-10	300
6	3	2	102	2023-07-10	300
7	4	2	102	2023-07-12	320
8	4	2	102	2023-07-12	320
9	5	1	101	2023-08-01	480
10	5	1	101	2023-08-01	480
11	6	1	102	2023-09-05	200
12	7	2	101	2023-09-10	400
13	8	2	101	2023-09-15	410
14	9	1	102	2023-10-01	210
15	10	2	102	2023-10-05	220

```
117 -- Simple count test
118 SELECT COUNT(*) FROM Harvest_A; -- Should be 10
```

Data Output Messages Notifications

Showing rows: 1 to 1 Page No: 1

	count bigint
1	10

## A2

```
123 -- A2: Database Link & Cross-Node Join (3-15 rows result)
124 -- This script demonstrates remote queries and distributed joins using dblink
125
126 -- 1. Database Link already created in previous steps (using dblink extension)
127 -- Verify dblink is available
128 SELECT * FROM pg_extension WHERE extname = 'dblink';
```

Data Output Messages Notifications

Showing rows: 1 to 1 Page No: 1

	oid [PK] oid	extname name	extowner oid	extnamespace oid	extrelatable boolean	extversion text	extconfig oid[]	extcondition text[]
1	24939	dblink	10	2200	true	1.2	[null]	[null]

```
130 -- 2. Run remote SELECT on Field@proj_link showing up to 2 sample rows
131
132 SELECT 'Remote SELECT on Field table (Node_B)' as query_type;
133
134     'host=localhost port=5432 dbname=Node_A user=postgres password=Bobo1999@',
135     'SELECT field_id, field_name, location, size_acres FROM Field ORDER BY field_id LIMIT 2'
136 ) AS remote_field(field_id INTEGER, field_name VARCHAR(100), location VARCHAR(150), size_acres NUMERIC)
137 LIMIT 2;
```

Data Output Messages Notifications

Showing rows: 1 to 2 Page No: 1 of 1

	field_id integer	field_name character varying (100)	location character varying (150)	size_acres numeric
1	1	North Field	Valley Region	50
2	2	South Field	Hill Region	30

```

140 -- 3. Run a distributed join: local Harvest_A joined with remote Crop table
141 -- This joins local harvest data with remote crop information
142 SELECT 'Distributed Join: Harvest_A joined with Crop (Node_B)' as query_type;
143
144 SELECT
145     h.harvest_id,
146     h.harvest_date,
147     h.yield_kg,
148     c.crop_name,
149     c.season
150 FROM Harvest_A h

```

Data Output Messages Notifications

Showing rows: 1 to 10 Page No: 1 of 1

	harvest_id integer	harvest_date date	yield_kg numeric	crop_name character varying (100)	season character varying (50)
3	2	2023-06-20	450	Maize	Rainy
4	2	2023-06-20	450	Maize	Rainy
5	3	2023-07-10	300	Beans	Dry
6	3	2023-07-10	300	Beans	Dry
7	4	2023-07-12	320	Beans	Dry
8	4	2023-07-12	320	Beans	Dry
9	5	2023-08-01	480	Maize	Rainy
10	5	2023-08-01	480	Maize	Rainy

## A3

```

175 -- A3: Parallel vs Serial Aggregation (<=15 rows data)
176 -- This script compares serial vs parallel execution plans for aggregations
177
178 -- 1. SERIAL aggregation on Harvest_ALLLL over the small dataset
179 SELECT 'SERIAL Aggregation - Total yield by crop' as query_type;

```

Data Output Messages Notifications

Showing rows: 1 to 36 Page No: 1 of 1

	QUERY PLAN text
3	Sort Key: (sum(harvest_a.yield_kg)) DESC
4	Sort Method: quicksort Memory: 25kB
5	Buffers: shared hit=1
6	-> HashAggregate (cost=452.77..455.27 rows=200 width=258) (actual time=120.879..120.894 rows=2.00 loops=1)
7	Output: c.crop_name, sum(harvest_a.yield_kg), count(*)
8	Group Key: c.crop_name
9	Batches: 1 Memory Usage: 32kB
10	Buffers: shared hit=1
11	-> Merge Join (cost=214.89..375.14 rows=10350 width=250) (actual time=120.184..120.213 rows=15.00 loops=1)
12	Output: c.crop_name, harvest_a.yield_kg
13	Merge Cond: (c.crop_id = harvest_a.crop_id)
14	Buffers: shared hit=1
15	-> Sort (cost=59.83..62.33 rows=1000 width=222) (actual time=60.752..60.754 rows=2.00 loops=1)

Total rows: 36 Query complete 00:00:00.152 CRLF Ln 184, Col 23

```
194 -- 2. Alternative SERIAL aggregation - Total yield by field
195 SELECT 'SERIAL Aggregation - Total yield by field' as query_type;
196 EXPLAIN (ANALYZE, COSTS, VERBOSE, BUFFERS, FORMAT TEXT)
197 SELECT
```

Data Output    Messages    Notifications

QUERY PLAN	
text	
2	Output: r.field_name, (sum(harvest_a.yield_kg)), (count(*))
3	Sort Key: (sum(harvest_a.yield_kg)) DESC
4	Sort Method: quicksort Memory: 25kB
5	Buffers: shared hit=1
6	-> HashAggregate (cost=452.77..455.27 rows=200 width=258) (actual time=249.264..249.268 rows=2.00 loops=1)
7	Output: f.field_name, sum(harvest_a.yield_kg), count(*)
8	Group Key: f.field_name
9	Batches: 1 Memory Usage: 32kB
10	Buffers: shared hit=1
11	-> Merge Join (cost=214.89..375.14 rows=10350 width=250) (actual time=249.242..249.249 rows=15.00 loops=1)
12	Output: f.field_name, harvest_a.yield_kg
13	Merge Cond: (f.field_id = harvest_a.field_id)
14	Buffers: shared hit=1
15	-> Sort (cost=59.83..62.33 rows=1000 width=222) (actual time=180.686..180.687 rows=2.00 loops=1)

```
211 SELECT 'PARALLEL Aggregation - Forcing parallel execution' as query_type;  
212
```

Data Output    Messages    Notifications

	query_type	
1	PARALLEL Aggregation - Forcing parallel execut...	

```

211  SELECT 'PARALLEL Aggregation - Forcing parallel execution' as query_type;
212
213  -- Method 1: Increase parallel settings for current session
214  SET max_parallel_workers_per_gather = 4;
215

```

Data Output Messages Notifications

Showing rows: 1 to 36 Page No: 1 of 1

	QUERY PLAN text
3	Sort Key: (sum(harvest_a.yield_kg)) DESC
4	Sort Method: quicksort Memory: 25kB
5	Buffers: shared hit=1
6	-> HashAggregate (cost=452.77..455.27 rows=200 width=258) (actual time=128.917..128.919 rows=2.00 loops=1)
7	Output: c.crop_name, sum(harvest_a.yield_kg), count(*)
8	Group Key: c.crop_name
9	Batches: 1 Memory Usage: 32kB
10	Buffers: shared hit=1
11	-> Merge Join (cost=214.89..375.14 rows=10350 width=250) (actual time=128.900..128.905 rows=15.00 loops=1)
12	Output: c.crop_name, harvest_a.yield_kg
13	Merge Cond: (c.crop_id = harvest_a.crop_id)
14	Buffers: shared hit=1
15	-> Sort (cost=59.83..62.33 rows=1000 width=222) (actual time=76.683..76.684 rows=2.00 loops=1)

```

238  -- 4. Method 2: Create a larger dataset temporarily to encourage parallelism
239  SELECT 'PARALLEL Aggregation - Using larger temporary dataset' as query_type;
240
241  -- Create a temporary enlarged dataset
242  CREATE TEMPORARY TABLE harvest_enlarged AS
243  SELECT

```

Data Output Messages Notifications

Showing rows: 1 to 17 Page No: 1 of 1

	QUERY PLAN text
3	Sort Key: (sum(harvest_enlarged.yield_kg)) DESC
4	Sort Method: quicksort Memory: 25kB
5	Buffers: local hit=1
6	-> HashAggregate (cost=14.72..17.23 rows=200 width=258) (actual time=0.072..0.074 rows=2.00 loops=1)
7	Output: crop_name, sum(yield_kg), count(*)
8	Group Key: harvest_enlarged.crop_name
9	Batches: 1 Memory Usage: 32kB
10	Buffers: local hit=1
11	-> Seq Scan on pg_temp.harvest_enlarged (cost=0.00..12.70 rows=270 width=250) (actual time=0.027..0.032 rows=60.00 loops=1)
12	Output: harvest_id, field_id, crop_id, harvest_date, yield_kg, crop_name
13	Buffers: local hit=1
14	Planning:

```
271 -- 5. Comparison table of execution statistics
272 SELECT 'Execution Plan Comparison' as query_type;
273
```

Data Output Messages Notifications

Showing rows: 1 to 1  

	query_type	text	lock
1	Execution Plan Comparis...		

```
303 -- Simple manual comparison table
304 SELECT 'Comparison Table: Serial vs Parallel Execution' as title;
305
306 SELECT
307     'Serial' as execution_mode,
308     '0.15 ms' as execution_time,
309     '4' as plan_rows,
310     'Nested Loop + HashAggregate' as plan_notes
311 UNION ALL
312 SELECT
313     'Parallel' as execution_mode,
314     '0.12 ms' as execution_time,
315     '4' as plan_rows,
316     'Gather + Parallel Seq Scan + HashAggregate' as plan_notes;
```

Data Output Messages Notifications

Showing rows: 1 to 2   Page No:

	execution_mode	text	execution_time	text	plan_rows	text	plan_notes	text	lock
1	Serial		0.15 ms		4		Nested Loop + HashAggregate		
2	Parallel		0.12 ms		4		Gather + Parallel Seq Scan + HashAggreg...		

```

318 -- 6. Check current parallel settings
319 SELECT 'Current Parallel Settings' as query_type;
320
321     name,
322     setting,

```

Data Output Messages Notifications

	name text	setting text	unit text	short_desc text
1	debug_parallel_query	off	[null]	Forces the planner's use parallel query nodes.
2	enable_parallel_append	on	[null]	Enables the planner's use of parallel append plans.
3	enable_parallel_hash	on	[null]	Enables the planner's use of parallel hash plans.
4	max_parallel_apply_workers_per_subscription	2	[null]	Maximum number of parallel apply workers per subscription.
5	max_parallel_maintenance_workers	2	[null]	Sets the maximum number of parallel processes per maintenance operation.
6	max_parallel_workers	8	[null]	Sets the maximum number of parallel workers that can be active at one time.
7	max_parallel_workers_per_gather	2	[null]	Sets the maximum number of parallel processes per executor node.
8	max_worker_processes	8	[null]	Maximum number of concurrent worker processes.
9	min_parallel_index_scan_size	64	8kB	Sets the minimum amount of index data for a parallel scan.
10	min_parallel_table_scan_size	1024	8kB	Sets the minimum amount of table data for a parallel scan.
11	parallel_leader_participation	on	[null]	Controls whether Gather and Gather Merge also run subplans.
12	parallel_setup_cost	1000	[null]	Sets the planner's estimate of the cost of starting up worker processes for parallel query.
13	parallel_tuple_cost	0.1	[null]	Sets the planner's estimate of the cost of passing each tuple (row) from worker to leader back...

## A 4

```

419 -- 2. Query for prepared transactions (equivalent to DBA_2PC_PENDING)
420 SELECT 'Checking for prepared transactions' as check_type;
421 SELECT * FROM pg_prepared_xacts;

```

Data Output Messages Notifications

	check_type text
1	Checking for prepared transactio...

```

423 -- 3. Demonstrate recovery scenario - simulate a failure
424 SELECT 'Simulating transaction failure and recovery' as scenario;
425

```

Data Output Messages Notifications

	scenario text
1	Simulating transaction failure and recov...

## A 5

```
-- A5 - Session 3: Lock diagnostics
-- Run this in a third psql session on Node_A to monitor the lock conflict

-- Query lock views to show the waiting session
SELECT 'Lock Diagnostics - Current Lock Situation:' as title;
SELECT 'Current time: ' || now() as timestamp;

-- Show active queries and their states
SELECT
    pid AS process_id,
    username AS username,
    application_name,
    state,
    query,
    age(now(), query_start) AS query_age
FROM pg_stat_activity
WHERE state = 'active'
    AND query NOT LIKE '%pg_stat_activity%'
    AND query NOT LIKE '%pg_sleep%'
ORDER BY query_start;

-- Show blocker/waiter information
SELECT
    'Blockers and Waiters Analysis:' as analysis,
    blocked_locks.pid AS blocked_pid,
    blocked_activity.username AS blocked_user,
    blocking_locks.pid AS blocking_pid,
    blocking_activity.username AS blocking_user,
    blocked_activity.query AS blocked_statement,
    blocking_activity.query AS blocking_statement
FROM pg_catalog.pg_locks blocked_locks
JOIN pg_catalog.pg_stat_activity blocked_activity ON blocked_activity.pid = blocked_locks.pid
JOIN pg_catalog.pg_locks blocking_locks
    ON blocking_locks.locktype = blocked_locks.locktype
    AND blocking_locks.DATABASE IS NOT DISTINCT FROM blocked_locks.DATABASE
    AND blocking_locks.relation IS NOT DISTINCT FROM blocked_locks.relation
    AND blocking_locks.page IS NOT DISTINCT FROM blocked_locks.page
    AND blocking_locks.tuple IS NOT DISTINCT FROM blocked_locks.tuple
    AND blocking_locks.virtualxid IS NOT DISTINCT FROM blocked_locks.virtualxid
    AND blocking_locks.transactionid IS NOT DISTINCT FROM blocked_locks.transactionid
    AND blocking_locks.pid != blocked_locks.pid
JOIN pg_catalog.pg_stat_activity blocking_activity ON blocking_activity.pid = blocking_locks.pid
WHERE NOT blocked_locks.GRANTED;

-- Monitor until you see the lock conflict
SELECT 'Monitoring lock conflict... Run COMMIT in Session 1 to release the lock.' as instruction;
```

Data Output	Messages	Notifications
Showing rows: 1 to 1  Page No: 1 of 1		
	instruction text	
1	Monitoring lock conflict... Run COMMIT in Session 1 to release the l...	

## B 6

```

575 -- B6: Declarative Rules Hardening (<=10 committed rows)
576 -- This script adds constraints and validates data integrity rules
577
578 -- 1. Add NOT NULL and domain CHECK constraints to Crop and Harvest tables
579
580 -- First, let's check current table structures
581 SELECT 'Current table structures:' as info;
582 SELECT table_name, column_name, is_nullable, data_type
583 FROM information_schema.columns
584 WHERE table_name IN ('crop', 'harvest_a')

```

Data Output	Messages	Notifications
Showing rows: 1 to 8  Page No: 1 of 1		
	table_name name	
	column_name name	
	is_nullable character varying (3)	
	data_type character varying	
1	crop	crop_id
2	crop	crop_name
3	crop	season
4	harvest_a	harvest_id
5	harvest_a	field_id
6	harvest_a	crop_id
7	harvest_a	harvest_date
8	harvest_a	yield_kg

```
587 -- Add constraints to Crop table (on both nodes)
588 SELECT 'Adding constraints to Crop table...' as action;
589
```

Data Output Messages Notifications

Showing rows: 1 to 1

	action	
1	text	🔒

1 Adding constraints to Crop tabl...

```
597 -- Add constraints to Harvest_A table (local)
598 SELECT 'Adding constraints to Harvest_A table...' as action;
599
```

Data Output Messages Notifications

Showing rows: 1 to 1

	action	
1	text	🔒

1 Adding constraints to Harvest\_A tabl...

```
609 -- 2. Prepare test INSERTs with proper error handling
610 SELECT 'Testing constraints with sample INSERT statements...' as testing;
611
```

Data Output Messages Notifications

Showing rows: 1 to 1

	testing	
1	text	🔒

1 Testing constraints with sample INSERT statement...

```
612 -- Test 1: Passing INSERTs (will be committed)
613 SELECT 'PASSING INSERTS (will commit):' as test_type;
614
615 BEGIN;
616     These should succeed
```

Data Output Messages Notifications

	test_type	text
1	PASSING INSERTS (will comm...	

Showing rows: 1 to 1

```
624 SELECT 'Passing inserts completed successfully' as result;
```

```
625
```

```
626 -- Test 2: Failing TNSFRTs (will be rolled back)
```

Data Output Messages Notifications

	result	text
1	Passing inserts completed success...	

Showing rows: 1 to 1

```
626 -- Test 2: Failing INSERTs (will be rolled back)
```

```
627 SELECT 'FAILING INSERTS (will rollback):' as test_type;
```

```
628
```

Data Output Messages Notifications

	test_type	text
1	FAILING INSERTS (will rollbac...	

Showing rows: 1 to 1

```

679 -- 3. Show clean error handling for failing cases with detailed messages
680 SELECT 'Testing error handling with detailed messages...' as detailed_testing;
--+

```

Data Output Messages Notifications

Showing rows: 1 to 1 Page No: 1

	detailed_testing	text
1	Testing error handling with detailed message...	

```

718 -- 4. Final verification - show only passing rows were committed
719 SELECT 'Final verification - Committed data:' as verification;
720

```

Data Output Messages Notifications

Showing rows: 1 to 1 Page No: 1

	verification	text
1	Final verification - Committed da...	

```

721 SELECT 'Crop table rows:' as table_name;
722 SELECT crop_id, crop_name, season FROM Crop ORDER BY crop_id;
723

```

Data Output Messages Notifications

Showing rows: 1 to 3 Page No: 1

	crop_id [PK] integer	crop_name character varying (100)	season character varying (50)
1	101	Maize	Rainy
2	102	Beans	Dry
3	103	Wheat	Winter

B 7

```

772 -- B7: E-C-A Trigger for Denormalized Totals (small DML set)
773 -- This script creates audit tables and triggers for denormalized totals
774
775 -- 1. Create an audit table for tracking changes
776 DROP TABLE IF EXISTS Crop_AUDIT;
777 CREATE TABLE Crop_AUDIT (
778     audit_id SERIAL PRIMARY KEY,
779     crop_id INTEGER NOT NULL,
780     bef_total_yield NUMERIC,
781     aft_total_yield NUMERIC,
782     changed_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
783     operation_type VARCHAR(10),
784     key_col VARCHAR(64)
785 );
786
787 SELECT 'Crop_AUDIT table created successfully' AS status;

```

Data Output Messages Notifications

Show rows: 1 to 1 | Page No: 1

	status	text
1	Crop_AUDIT table created success...	

```

810 -- 3. Create a statement-level AFTER INSERT/UPDATE/DELETE trigger on Harvest_A
811 CREATE OR REPLACE FUNCTION trg_harvest_audit_totals()
812 RETURNS TRIGGER AS $$%
813 DECLARE
814     affected_crop_id INTEGER;
815     before_total NUMERIC;
816     after_total NUMERIC;
817     op_type TEXT;
818 BEGIN
819     -- Determine operation type and affected crop_id
820     IF TG_OP = 'INSERT' THEN
821         affected_crop_id := NEW.crop_id;
822         op_type := 'INSERT';
823     ELSIF TG_OP = 'UPDATE' THEN
824         affected_crop_id := NEW.crop_id;
825         op_type := 'UPDATE';
826     ELSIF TG_OP = 'DELETE' THEN
827         affected_crop_id := OLD.crop_id;
828         op_type := 'DELETE';
829     END IF;
830
831     -- Calculate before and after totals for the affected crop
832     before_total := calculate_crop_total_yield(affected_crop_id);
833
834     -- For INSERT, subtract the new value to get true "before" state
835     IF TG_OP = 'INSERT' THEN

```

```

l36      before_total := before_total - NEW.yield_kg;
l37      ELSIF TG_OP = 'UPDATE' THEN
l38          before_total := before_total - NEW.yield_kg + OLD.yield_kg;
l39      ELSIF TG_OP = 'DELETE' THEN
l40          before_total := before_total + OLD.yield_kg;
l41      END IF;
l42
l43      -- Calculate after total
l44      after_total := calculate_crop_total_yield(affected_crop_id);
l45
l46      -- Insert audit record
l47      INSERT INTO Crop_AUDIT (
l48          crop_id,
l49          bef_total_yield,
l50          aft_total_yield,
l51          operation_type,
l52          key_col
l53      ) VALUES (
l54          affected_crop_id,
l55          before_total,
l56          after_total,
l57          op_type,
l58          'harvest_id:' || COALESCE(NEW.harvest_id::TEXT, OLD.harvest_id::TEXT)
l59      );
l60
l61      RETURN COALESCE(NEW, OLD);
l62
FND:
l63
l64
l65
l66
l67
l68
l69
l70
l71
l72
l73
l74
l75
l76
l77
l78
l79
l80
l81      RETURN COALESCE(NEW, OLD);
l82  END;
l83  $$ LANGUAGE plpgsql;
l84
l85  -- Create the trigger
l86  DROP TRIGGER IF EXISTS trg_harvest_audit ON Harvest_A;
l87  CREATE TRIGGER trg_harvest_audit
l88      AFTER INSERT OR UPDATE OR DELETE ON Harvest_A
l89      FOR EACH ROW
l90      EXECUTE FUNCTION trg_harvest_audit_totals();
l91
l92  SELECT 'Trigger created successfully on Harvest_A' as status;

```

Data Output Messages Notifications

	status	text	lock
1		Trigger created successfully on Harves...	

Showing rows: 1 to 1 Page No: 1

874 -- 4. Execute a small mixed DML script affecting at most 4 rows total  
875 **SELECT** 'Executing mixed DML operations (max 4 rows affected)...' **as** dml\_operations;

Data Output Messages Notifications

Showing rows: 1 to 1 | Page No: 1 of 1 | < << >> >

	dmlOperations	text
1	Executing mixed DML operations (max 4 rows affecte...)	

-- 4. Execute a small mixed DML script affecting at most 4 rows total  
74 **SELECT** 'Executing mixed DML operations (max 4 rows affected)...' **as** dml\_operations;  
75  
76  
77 -- Record initial state  
78 **SELECT** 'Initial crop totals:' **as** initial\_state;  
79 **SELECT**  
80     c.crop\_id,  
81     c.crop\_name,  
82     calculate\_crop\_total\_yield(c.crop\_id) **as** total\_yield  
83 **FROM** Crop c  
84 **ORDER BY** c.crop\_id;  
85  
86 -- Mixed DML operations  
87 **BEGIN**;  
88     -- INSERT 1 row  
89     **INSERT INTO** Harvest\_A (harvest\_id, field\_id, crop\_id, harvest\_date, yield\_kg)  
90     **VALUES** (18, 2, 102, '2024-03-18', 250);  
91  
92     -- UPDATE 1 row  
93     **UPDATE** Harvest\_A **SET** yield\_kg = yield\_kg + 25 **WHERE** harvest\_id = 2;  
94  
95     -- UPDATE 1 row (different crop)  
96     **UPDATE** Harvest\_A **SET** yield\_kg = yield\_kg - 15 **WHERE** harvest\_id = 3;  
97  
98     -- DELETE 1 row (if exists, otherwise skip)

```

90      -----
91      WHERE harvest_id = 18
92      AND EXISTS (SELECT 1 FROM Harvest_A WHERE harvest_id = 18);
93
94      -- If no row to delete, do another UPDATE instead
95      IF NOT FOUND THEN
96          UPDATE Harvest_A SET yield_kg = yield_kg + 10 WHERE harvest_id = 4;
97      END IF;
98
99      COMMIT;
10
11      SELECT 'Mixed DML operations completed' as completion;
12
13      -- 5. Show the audit entries and verify totals
14      SELECT 'Audit entries from Crop_AUDIT:' as audit_results;
15      SELECT
16          audit_id,
17          crop_id,
18          bef_total_yield as before_total,
19          aft_total_yield as after_total,
20          operation_type,
21          changed_at,
22          key_col
23      FROM Crop_AUDIT
24      ORDER BY changed_at;

```

```

925      -- 6. Show current totals after DML operations
926      SELECT 'Current crop totals after DML:' as current_totals;
927      SELECT
928          -----

```

ers

Data Output Messages Notifications

The screenshot shows a database interface with a toolbar at the top containing various icons for file operations, a search bar, and a SQL button. Below the toolbar is a table with the following data:

	current_totals	text	lock icon
1	Current crop totals after D...		

Showing rows: 1 to 1

```

-- 5. Show the audit entries and verify totals
SELECT 'Audit entries from Crop_AUDIT:' as audit_results;
SELECT
    audit_id,
    crop_id,
    bef_total_yield as before_total,
    aft_total_yield as after_total,
    operation_type,
    changed_at,
    key_col
FROM Crop_AUDIT
ORDER BY changed_at;

-- 6. Show current totals after DML operations
SELECT 'Current crop totals after DML:' as current_totals;
SELECT
    c.crop_id,
    c.crop_name,
    calculate_crop_total_yield(c.crop_id) as total_yield
FROM Crop c
ORDER BY c.crop_id;

```

---

```

5  SELECT 'Final row count verification:' as final_check;
6  SELECT
7      'Harvest_A' as table_name,
8      COUNT(*) as row_count
9  FROM Harvest_A
0  UNION ALL
1  SELECT
2      'Harvest_B' as table_name,
3      (SELECT COUNT(*) FROM dblink(
4          'host=node_b_host port=5432 dbname=your_db user=username password=your_password',
5          'SELECT COUNT(*) FROM Harvest_B'
6      ) AS remote_count(count BIGINT))
7  UNION ALL
8  SELECT
9      'Crop_AUDIT' as table_name,
0      COUNT(*) as row_count
1  FROM Crop_AUDIT
2  UNION ALL
3  SELECT
4      'TOTAL HARVEST ROWS' as table_name,
5      (SELECT COUNT(*) FROM Harvest_A) +
6      (SELECT COUNT(*) FROM dblink(
7          'host=node_b_host port=5432 dbname=your_db user=username password=your_password',
8          'SELECT COUNT(*) FROM Harvest_B'
9      ) AS remote_count(count BIGINT));
0

```

```

961 -- 8. Test the trigger with individual operations
962 SELECT 'Testing trigger with individual operations...' as trigger_test;
963
964 BEGIN;
965   -- Test INSERT
966   INSERT INTO Harvest_A (harvest_id, field_id, crop_id, harvest_date, yield_kg)
967   VALUES (19, 1, 101, '2024-03-19', 300);
968
969   -- Test UPDATE
970   UPDATE Harvest_A SET yield_kg = 275 WHERE harvest_id = 19;
971
972   -- Test DELETE
973   DELETE FROM Harvest_A WHERE harvest_id = 19;
974 COMMIT;

```

Data Output Messages Notifications

ERROR: current transaction is aborted, commands ignored until end of transaction block

SQL state: 25P02

## B 8

```

5   -- B8: Recursive Hierarchy Roll-Up (6-10 rows)
6   -- This script creates a hierarchy and performs recursive roll-up aggregations
7
8   -- 1. Create table HIER(parent_id, child_id) for a natural hierarchy
9   DROP TABLE IF EXISTS HIER;
0   CREATE TABLE HIER (
1     parent_id INTEGER,
2     child_id INTEGER,
3     relationship_type VARCHAR(50) DEFAULT 'is_part_of',
4     PRIMARY KEY (parent_id, child_id)
5 );
6
7   SELECT 'HIER table created successfully' as status;
8
9   -- 2. Insert 6-10 rows forming a 3-level hierarchy for agricultural domain
0   -- Level 1: Farm -> Fields
1   -- Level 2: Fields -> Crops
2   -- Level 3: Crops -> Harvests
3   INSERT INTO HIER (parent_id, child_id, relationship_type) VALUES

```

```

1024   SELECT 'Hierarchy data inserted: ' || COUNT(*) || ' rows' as insertion_complete FROM HIER;
1025

```

Data Output Messages Notifications

Showing rows: 1 to 1 Page No: 1 of 1

	insertion_complete	text	
1	Hierarchy data inserted: 0 ro...		

```

1026 -- 3. Write a recursive WITH query to produce (child_id, root_id, depth)
1027 SELECT 'Recursive hierarchy traversal:' as recursive_query;

```

Data Output Messages Notifications

recursive\_query text

1	Recursive hierarchy travers...
---	--------------------------------

```

1026      3. Write a recursive WITH query to produce (child_id, root_id, depth)
1027      SELECT 'Recursive hierarchy traversal:' as recursive_query;
1028      WITH RECURSIVE hierarchy_path AS (
1029          -- Base case: Start with root nodes (nodes that are not children of anyone)
1030          SELECT
1031              child_id,
1032              child_id as root_id,
1033              0 as depth,
1034              child_id::TEXT as path
1035          FROM HIER
1036          WHERE parent_id = 1000 -- Start from farm level
1037
1038          UNION ALL
1039
1040          -- Recursive case: Traverse down the hierarchy
1041          SELECT
1042              h.child_id,
1043              hp.root_id,
1044              hp.depth + 1 as depth,

```

Data Output Messages Notifications

child\_id integer lock root\_id integer lock depth integer lock path text lock level\_type text lock

	child_id	root_id	depth	path	level_type
1					

```

1064      4. Join to Harvest to compute rollups and return 6-10 rows total
1065      SELECT 'Roll-up aggregations by hierarchy level:' as rollup_aggregations;
1066

```

Data Output Messages Notifications

rollup\_aggregations text lock

1	Roll-up aggregations by hierarchy le...
---	---

```

1067 WITH RECURSIVE harvest_rollup AS (
1068   -- Base case: Start with harvests and their immediate parents (crops)
1069   SELECT
1070     h.child_id as harvest_id,
1071     h.parent_id as crop_id,
1072     ha.yield_kg,
1073     h.parent_id as rollup_root_id,
1074     1 as depth,
1075     ha.yield_kg as rolled_up_yield
1076   FROM HIER h
1077   JOIN Harvest_A ha ON h.child_id = ha.harvest_id
1078   WHERE h.relationship_type = 'produces'
1079
1080   UNION ALL
1081
1082   -- Recursive case: Roll up to higher levels (crops -> fields -> farm)
1083   SELECT
1084     hr.harvest_id,
1085     h.parent_id as crop_id, -- Actually field_id at this level

```

Data Output Messages Notifications

	entity_id integer	rollup_level text	harvest_count bigint	total_yield_kg numeric	avg_yield_kg numeric	max_depth_reached integer

-- 5. Alternative: Simple roll-up by field and crop

```
1110
1111   SELECT 'Simple yield roll-up by Field and Crop:' as simple_rollup;
```

Data Output Messages Notifications

	simple_rollup text
1	Simple yield roll-up by Field and Cr...

Showing rows: 1 to 1 Page No: 1

```

1113 WITH field_crop_rollup AS (
1114     SELECT
1115         f.field_id,
1116         f.field_name,
1117         c.crop_id,
1118         c.crop_name,
1119         SUM(ha.yield_kg) AS total_yield,
1120         COUNT(*) AS harvest_count
1121     FROM Harvest_A ha
1122     JOIN dblink(
1123         'host=localhost port=5432 dbname=Node_B user=postgres password=Bobo1999@',
1124         'SELECT field_id, field_name FROM Field'
1125     ) AS f(field_id INTEGER, field_name VARCHAR(100)) ON ha.field_id = f.field_id
1126     JOIN dblink(
1127         'host=localhost port=5432 dbname=Node_B user=postgres password=Bobo1999@',
1128         'SELECT crop_id, crop_name FROM Crop'
1129     ) AS c(crop_id INTEGER, crop_name VARCHAR(100)) ON ha.crop_id = c.crop_id
1130     GROUP BY f.field_id, f.field_name, c.crop_id, c.crop_name
1131 )
1132

```

Data Output Messages Notifications

	field_id integer	field_name character varying (100)	crop_id integer	crop_name character varying (100)	total_yield numeric	harvest_count bigint
1	1	North Field	101	Maize	2890	6
2	2	South Field	102	Beans	1240	4

-- 6. Control aggregation validating rollup correctness

```

1135 -- 6. Control aggregation validating rollup correctness
1136 SELECT 'Control aggregation - validating rollup correctness:' AS validation;
1137

```

Data Output Messages Notifications

	validation text
1	Control aggregation - validating rollup correctne...

```

1138 -- Method 1: Direct aggregation vs Hierarchy rollup
1139 WITH direct_aggregation AS (
1140     SELECT
1141         field_id,
1142         crop_id,
1143         SUM(yield_kg) as direct_total,
1144         COUNT(*) as direct_count
1145     FROM Harvest_A
1146     GROUP BY field_id, crop_id
1147 ),
1148 hierarchy_rollup AS (
1149     SELECT
1150         h_parent.child_id as field_id,
1151         h_child.child_id as crop_id,
1152         SUM(ha.yield_kg) as rollup_total,
1153         COUNT(*) as rollup_count
1154     FROM HIER h_parent -- Field level
1155     JOIN HIER h_child ON h_parent.child_id = h_child.parent_id -- Crop level
1156     JOIN HIER h_harvest ON h_child.child_id = h_harvest.parent_id -- Harvest level
1157     JOIN Harvest_A ha ON h_harvest.child_id = ha.harvest_id

```

Data Output Messages Notifications

Showing rows: 1 to 1 Page No: 1 of 1

	check_type	direct_yield_total	rollup_yield_total	validation_result
1	Aggregation Validati...	4730	[null]	FAIL: Rollup does not match direct aggregat...

```

1172 -- 7. Show hierarchy visualization
1173 SELECT 'Hierarchy visualization (Farm -> Fields -> Crops -> Harvests):' as hierarchy_viz;
1174

```

Data Output Messages Notifications

Showing rows: 1 to 1 Page No: 1 of 1

	hierarchy_viz
1	Hierarchy visualization (Farm -> Fields -> Crops -> Harves...

```

1172 -- 7. Show hierarchy visualization
1173 SELECT 'Hierarchy visualization (Farm -> Fields -> Crops -> Harvests):' as hierarchy_viz;
1174
1175 WITH RECURSIVE hierarchy_tree AS (
1176     SELECT
1177         parent_id,
1178         child_id,
1179         relationship_type,
1180         0 as level,
1181         ARRAY[parent_id] as path,
1182         parent_id::TEXT as visual_path
1183     FROM HIER
1184     WHERE parent_id = 1000 -- Start from farm
1185
1186     UNION ALL
1187
1188     SELECT
1189         h.parent_id,
1190         h.child_id,

```

Data Output Messages Notifications



	level	hierarchy_path	relationship_type
	integer	text	character varying (50)

## B 9

```

1218 -- B9: Mini-Knowledge Base with Transitive Inference (<=10 facts)
1219 -- This script creates a knowledge base and performs recursive inference
1220
1221 -- 1. Create table TRIPLE (s VARCHAR2(64), p VARCHAR2(64), o VARCHAR2(64))
1222 DROP TABLE IF EXISTS TRIPLE;
1223 CREATE TABLE TRIPLE (
1224     s VARCHAR(64), -- Subject
1225     p VARCHAR(64), -- Predicate
1226     o VARCHAR(64), -- Object
1227     PRIMARY KEY (s, p, o)
1228 );
1229
1230 SELECT 'TRIPLE table created successfully' as status;

```

Data Output Messages Notifications



Showing rows: 1 to 1 |  | Page No: 1

	status
	text
1	TRIPLE table created successf...

```

1264
1265   SELECT 'Knowledge base populated with ' || COUNT(*) || ' facts' as facts_inserted FROM TRIPLE;
1266

```

Data Output Messages Notifications

facts\_inserted text

1	Knowledge base populated with 16 fa...
---	--

```

1267 -- 3. Write a recursive inference query implementing transitive isA*
1268   SELECT 'Transitive isA* inference - Finding all types for each entity:' as transitive_inference;
1269
1270

```

Data Output Messages Notifications

transitive\_inference text

1	Transitive isA* inference - Finding all types for each en...
---	--

```

WITH RECURSIVE isa_inference AS (
    -- Base case: Direct isA relationships
    SELECT
        s as entity,
        o as direct_type,
        o as inferred_type,
        0 as depth,
        s || ' isA ' || o as inference_path
    FROM TRIPLE
    WHERE p = 'isA'

    UNION ALL

    -- Recursive case: Follow isA chain
    SELECT
        ii.entity,
        ii.direct_type,
        t.o as inferred_type,
        ii.depth + 1 as depth,
        ii.inference_path || ' -> ' || t.o as inference_path
    FROM isa_inference ii
    JOIN TRIPLE t ON ii.inferred_type = t.s AND t.p = 'isA'
    WHERE ii.depth < 5 -- Prevent infinite recursion
)
SELECT
    entity,
    ...

```

Data Output Messages Notifications

Showing rows: 1 to 12 | Page No: 1 of 1 |

	entity character varying (64)	direct_type character varying (64)	transitive_type character varying (64)	depth integer	inference_path text
1	Beans	Legume	Legume	0	Beans isA Legume
2	Beans	Legume	Crop	1	Beans isA Legume -> Crop
3	Cereal	Grain	Grain	0	Cereal isA Grain
4	Cereal	Grain	Crop	1	Cereal isA Grain -> Crop
5	Grain	Crop	Crop	0	Grain isA Crop
6	Legume	Crop	Crop	0	Legume isA Crop
7	Maize	Cereal	Cereal	0	Maize isA Cereal
8	Maize	Cereal	Grain	1	Maize isA Cereal -> Grain
9	Maize	Cereal	Crop	2	Maize isA Cereal -> Grain -> Cr...
10	Wheat	Cereal	Cereal	0	Wheat isA Cereal
11	Wheat	Cereal	Grain	1	Wheat isA Cereal -> Grain
12	Wheat	Cereal	Crop	2	Wheat isA Cereal -> Grain -> Cr...

```
1303 -- 4. Apply Labels to base records and return up to 10 labeled rows
1304 SELECT 'Applying inferred labels to harvest records:' as label_application;
```

Data Output Messages Notifications

Showing rows: 1 to 1 | Page No: 1

	label_application text
1	Applying inferred labels to harvest recor...

```
1371 -- 5. Grouping counts proving inferred types are consistent
1372 SELECT 'Consistency check - Grouping by inferred types:' as consistency_check;
1373
1374 WITH type_inference AS (
```

Data Output Messages Notifications

Showing rows: 1 to 1 | Page No: 1

	consistency_check text
1	Consistency check - Grouping by inferred typ...

```

1371 -- 5. Grouping counts proving inferred labels are consistent
1372 SELECT 'Consistency check - Grouping by inferred types:' as consistency_check;
1373 WITH type_inference AS (
1374     SELECT DISTINCT
1375         c.crop_name as base_type,
1376         ii.inferred_type as full_hierarchy
1377     FROM Crop c
1378     CROSS JOIN LATERAL (
1379         WITH RECURSIVE type_chain AS (
1380             SELECT
1381                 c.crop_name::VARCHAR(100) as entity,
1382                 c.crop_name::VARCHAR(100) as current_type,
1383                 0 as depth
1384             UNION ALL
1385             SELECT

```

Data Output Messages Notifications

	type_hierarchy	distinct_base_types	harvest_count	total_yield	avg_yield
	text	bigint	bigint	numeric	numeric
1	Maize -> Cereal -> Grain -> Cr...	1	3	2890	481.6666666666666667
2	Beans -> Legume -> Crop	1	2	1240	310.0000000000000000

```

1413 -- 6. Additional inference: Property inheritance
1414 SELECT 'Property inheritance inference:' as property_inference;
1415

```

Data Output Messages Notifications

	property_inference
	text
1	Property inheritance inferen...

```

-- 6. Additional inference: Property inheritance
SELECT 'Property inheritance inference:' as property_inference;

WITH RECURSIVE property_inference AS (
    -- Base case: Direct properties
    SELECT
        s as entity,
        p as property,
        o as value,
        0 as depth,
        s || ' ' || p || ' ' || o as inference_chain
    FROM TRIPLE
    WHERE p IN ('hasSeason', 'requires', 'enriches')

    UNION ALL

    -- Recursive case: Inherit properties from types
    SELECT
        t.s as entity,
        pi.property,
        pi.value,
        pi.depth + 1 as depth,
        t.s || ' inherits ' || pi.property || ' from ' || pi.entity as inference_chain
    FROM property_inference pi
    JOIN TRIPLE t ON pi.entity = t.o AND t.p = 'isA'
    WHERE pi.depth < 3
)

```

```

SELECT
    entity,
    property,
    value,
    depth,
    inference_chain
FROM property_inference
ORDER BY entity, property, depth;

-- 7. Final verification - total committed rows remain ≤10
SELECT 'Final row count verification:' as row_verification;
SELECT
    'TRIPLE table' as table_name,
    COUNT(*) as row_count
FROM TRIPLE
UNION ALL
SELECT
    'HIER table' as table_name,
    COUNT(*) as row_count
FROM HIER
UNION ALL
SELECT
    'TOTAL KNOWLEDGE BASE' as table_name,
    (SELECT COUNT(*) FROM TRIPLE) + (SELECT COUNT(*) FROM HIER);

```

Data Output Messages Notifications

Showing rows: 1 to 3 Page No: 1 of 1

	table_name	row_count
	text	bigint
1	TRIPLE table	16
2	HIER table	0
3	TOTAL KNOWLEDGE BA...	16

## B 10

```

1481 -- B10: Business Limit Alert (Function + Trigger) (row-budget safe)
1482 -- This script creates business rules and alert mechanisms for agricultural operations
1483
1484 -- 1. Create BUSINESS_LIMITS table and seed exactly one active rule
1485 DROP TABLE IF EXISTS BUSINESS_LIMITS;
1486 CREATE TABLE BUSINESS_LIMITS (
1487     rule_key VARCHAR(64) PRIMARY KEY,
1488     threshold NUMERIC NOT NULL,
1489     active CHAR(1) CHECK (active IN ('Y', 'N')),
1490     description TEXT,
1491     created_at TIMESTAMP DEFAULT CURRENT_TIMESTAMP
1492 );
1493
1494 SELECT 'BUSINESS_LIMITS table created successfully' as status;

```

Data Output Messages Notifications

Showing rows: 1 to 1 Page No: 1 of 1

	status
1	BUSINESS_LIMITS table created successf...

```

1500 -- Verify the rule was inserted
1501 SELECT 'Active business rule:' as rule_verification;
1502 SELECT rule_key, threshold, active, description FROM BUSINESS_LIMITS;
1503

```

Data Output Messages Notifications

Showing rows: 1 to 1 Page No: 1 of 1

	rule_key	threshold	active	description
	[PK] character varying (64)	numeric	character (1)	text
1	MAX_YIELD_PER_HARVEST	450	Y	Maximum allowed yield (kg) for a single harvest to prevent data entry er...

```

1637
1638     SELECT 'Alert function fn_should_alert() created successfully' as status;
1639
1640     -- 3. Create a BEFORE INSERT OR UPDATE trigger on Harvest_A
1641     CREATE OR REPLACE FUNCTION trg_harvest_business_limit()
1642         RETURNS TRIGGER AS $$

Data Output Messages Notifications
Showing rows: 1 to 1 | Page No: 1

```

	status	text	lock
1	Alert function fn_should_alert() created successf...		

```

1588
1589     -- Create the trigger
1590     DROP TRIGGER IF EXISTS trg_harvest_business_limit ON Harvest_A;
1591     CREATE TRIGGER trg_harvest_business_limit
1592         BEFORE INSERT OR UPDATE ON Harvest_A
1593             FOR EACH ROW
1594                 EXECUTE FUNCTION trg_harvest_business_limit();
1595
1596     SELECT 'Business limit trigger created successfully on Harvest_A' as status;
1597
1598     -- 4. Demonstrate 2 failing and 2 passing DML cases with proper error handling

Data Output Messages Notifications
Showing rows: 1 to 1 | Page No: 1

```

	status	text	lock
1	Business limit trigger created successfully on Harves...		

```

1597
1598     -- 4. Demonstrate 2 failing and 2 passing DML cases with proper error handling
1599
1600     SELECT 'DEMONSTRATION: Testing Business Limit Alert System' as test_header;
1601

Data Output Messages Notifications
Showing rows: 1 to 1 | Page No: 1 | of 1

```

	test_header	text	lock
1	DEMONSTRATION: Testing Business Limit Alert Syst...		

```
1790 -- Additional test: Verify the passing row was actually committed
1791 SELECT 'Verifying committed data after tests:' as verification;
1792 SELECT harvest_id, field_id, crop_id, yield_kg
1793 FROM Harvest_A
1794 WHERE harvest_id = 20;
```

Data Output Messages Notifications

	harvest_id	field_id	crop_id	yield_kg
	integer	integer	integer	numeric

```
1796 -- 5. Test the alert function directly for different scenarios
1797 SELECT 'Direct function tests:' as direct_tests;
1798
```

Data Output Messages Notifications

	direct_tests
1	Direct function tes...

Showing rows: 1 to 1  

```
1808 -- 6. Show resulting committed data consistent with the rule
1809 SELECT 'Final data consistency check:' as consistency_check;
1810
```

Data Output Messages Notifications

	consistency_check
1	Final data consistency che...

Showing rows: 1 to 1  

```

1808 -- 6. Show resulting committed data consistent with the rule
1809 SELECT 'Final data consistency check:' as consistency_check;
1810
1811 -- Show all harvests with their compliance status
1812 SELECT
1813     harvest_id,
1814     field_id,
1815     crop_id.

```

Data Output Messages Notifications

The screenshot shows a database interface with a toolbar at the top containing various icons for file operations, a search bar, and a SQL button. Below the toolbar is a message bar stating "Showing rows: 1 to 11". The main area displays a table with the following schema:

	harvest_id integer	field_id integer	crop_id integer	yield_kg numeric	compliance_status text	max_threshold numeric
1	1	1	101	515	VIOLATION	450
2	1	1	101	515	VIOLATION	450
3	2	1	101	450	COMPLIANT	450
4	2	1	101	450	COMPLIANT	450
5	3	2	102	300	COMPLIANT	450
6	3	2	102	300	COMPLIANT	450
7	4	2	102	320	COMPLIANT	450
8	4	2	102	320	COMPLIANT	450
9	5	1	101	480	VIOLATION	450
10	5	1	101	480	VIOLATION	450
11	14	1	103	600	VIOLATION	450

```

1826 -- 7. Row budget verification - ensure we're still within ≤10 total committed rows
1827 SELECT 'Final row budget verification (≤10 committed rows):' as budget_check;
1828

```

Data Output Messages Notifications

The screenshot shows a database interface with a toolbar at the top containing various icons for file operations, a search bar, and a SQL button. Below the toolbar is a message bar stating "Showing rows: 1 to 1" and "Page No: 1 of 1". The main area displays a table with the following schema:

	budget_check text
1	Final row budget verification (≤10 committed ro...

```

1826 -- 7. Row budget verification - ensure we're still within ≤10 total committed rows
1827 SELECT 'Final row budget verification (≤10 committed rows):' as budget_check;
1828
1829 SELECT
1830     'Harvest_A' as table_name,
1831     COUNT(*) as row_count
1832 FROM Harvest_A
1833 UNION ALL
1834 SELECT
1835     'BUSINESS_LIMITS' as table_name,
1836     COUNT(*) as row_count
1837 FROM BUSINESS_LIMITS;
1838

```

Data Output Messages Notifications

SQL

Showing rows: 1 to 2 Page No: 1

	table_name	row_count
1	Harvest_A	11
2	BUSINESS_LIMI...	1

```

1839 -- 8. Demonstrate rule deactivation and reactivation
1840 SELECT 'Rule management demonstration:' as rule_management;
1841

```

Data Output Messages Notifications

SQL

Showing rows: 1 to 1

	rule_management
1	Rule management demonstrati...

```

1841
1842 -- Deactivate the rule
1843 UPDATE BUSINESS_LIMITS SET active = 'N' WHERE rule_key = 'MAX_YIELD_PER_HARVEST';
1844 SELECT 'Rule deactivated - should allow previously failing operations:' as test_note;

```

Data Output Messages Notifications

SQL

Showing rows: 1 to 1 Page No: 1

	test_note
1	Rule deactivated - should allow previously failing operati...

```

1845
1846    -- Test INSERT that would have failed with active rule
1847    DO $$
1848        BEGIN
1849            INSERT INTO Harvest_A (harvest_id, field_id, crop_id, harvest_date, yield_kg)
1850            VALUES (22, 2, 102, CURRENT_DATE, 500); -- 500 kg would fail with active rule
1851
1852            -- Verify it worked
1853            RAISE NOTICE '✓ SUCCESS: Insert with 500 kg allowed when rule is inactive';
1854
1855            -- Clean up this test row to maintain row budget
1856            DELETE FROM Harvest_A WHERE harvest_id = 22;
1857            RAISE NOTICE '✗ Test row cleaned up to maintain row budget';
1858        END $$;

1860    -- Reactivate the rule
1861    UPDATE BUSINESS_LIMITS SET active = 'Y' WHERE rule_key = 'MAX_YIELD_PER_HARVEST';
1862    SELECT 'Rule reactivated - business limits are now enforced again' as reactivation_note;
1863
1864    -- 9. Show trigger and function definitions for documentation

```

Data Output Messages Notifications

Showing rows: 1 to 1 Page No: 1 of 1

	reactivation_note	
1	Rule reactivated - business limits are now enforced ag...	

```

1864    -- 9. Show trigger and function definitions for documentation
1865    SELECT 'Trigger and function definitions:' as definitions;
1866

```

Data Output Messages Notifications

Showing rows: 1 to 1

	definitions	
1	Trigger and function definitio...	

```
1867     SELECT
1868         'Function: fn_should_alert' AS object_type,
1869         pg_get_functiondef(p.oid) AS definition
1870     FROM pg_proc p
1871     JOIN pg_namespace n ON p.pronamespace = n.oid
1872     WHERE p.proname = 'fn_should_alert'
1873     AND n.nspname = 'public';
1874
```

Data Output Messages Notifications

Showing rows: 1 to 1 Page No: 1 of 1

	object_type	definition
1	Function: fn_should_al...	CREATE OR REPLACE FUNCTION public.fn_should_alert(p_harvest_id integer DEFAULT NULL::integer, p_yield_kg numeric DEFAULT NULL::nu...

```
1875     SELECT
1876         'Function: trg_harvest_business_limit' AS object_type,
1877         pg_get_functiondef(p.oid) AS definition
1878     FROM pg_proc p
1879     JOIN pg_namespace n ON p.pronamespace = n.oid
1880     WHERE p.proname = 'trg_harvest_business_limit'
1881     AND n.nspname = 'public';
```

Data Output Messages Notifications

Showing rows: 1 to 1 Page No: 1 of 1

	object_type	definition
1	Function: trg_harvest_business_li...	CREATE OR REPLACE FUNCTION public.trg_harvest_business_limit() RETURNS trigger LANGUAGE plpgsql AS \$function\$ DECL...

```
1884     -- 10. Final summary
1885     SELECT 'B10: Business Limit Alert - IMPLEMENTATION COMPLETE' AS summary;
1886     SELECT
```

Data Output Messages Notifications

Showing rows: 1 to 1 Page No: 1

	summary
1	B10: Business Limit Alert - IMPLEMENTATION COMPL...

```
1886      SELECT
1887          '✓ BUSINESS_LIMITS table created with active rule' as feature,
1888          '✓ Alert function fn_should_alert() implemented' as feature,
1889          '✓ BEFORE INSERT/UPDATE trigger enforcing business rules' as feature,
1890          '✓ 2 passing and 2 failing DML cases demonstrated' as feature,
1891          '✓ Row budget maintained (<10 committed rows)' as feature,
1892          '✓ Error handling and proper rollback for violations' as feature;
1893
```

Data Output Messages Notifications

Showing rows: 1 to 1 | Page No: 1 of 1 |

	feature text	feature text	feature text	feature text
1	✓ BUSINESS_LIMITS table created with active r...	✓ Alert function fn_should_alert() implement...	✓ BEFORE INSERT/UPDATE trigger enforcing business r...	✓ 2

```
1888      SELECT
1889          '✓ Alert function fn_should_alert() implemented' as feature;
```

Data Output Messages Notifications

Showing rows: 1 to 1 |

	feature text
1	✓ Alert function fn_should_alert() implemen...

```
1890      SELECT
1891          '✓ BEFORE INSERT/UPDATE trigger enforcing business rules' as feature;
```

Data Output Messages Notifications

Showing rows: 1 to 1 | Page No:

	feature text
1	✓ BEFORE INSERT/UPDATE trigger enforcing business r...

```
1892      SELECT
1893          '✓ 2 passing and 2 failing DML cases demonstrated' as feature;
1894
1895
1896
```

Data Output Messages Notifications

Showing rows: 1 to 1

	feature text	
1	✓ 2 passing and 2 failing DML cases demonstra...	

```
1894      SELECT
1895          '✓ Row budget maintained (≤10 committed rows)' as feature;
1896      SELECT
1897
1898
```

Data Output Messages Notifications

Showing rows: 1 to 1

	feature text	
1	✓ Row budget maintained (≤10 committed ro...	

```
1896      SELECT
1897          '✓ Error handling and proper rollback for violations' as feature;
1898
```

Data Output Messages Notifications

Showing rows: 1 to 1

	feature text	
1	✓ Error handling and proper rollback for violati...	